Mathew Stevens

mathewstevens7457@gmail.com | https://mathewstevens.dev/portfolio/ | https://github.com/mstevens843
Honeycomb.io - Senior Software Engineer - LLM Observability

Dear Honeycomb hiring team,

I am excited about the Senior Software Engineer - LLM Observability role because it is focused on one of the most important problems in production AI: making agent and LLM systems observable enough that engineers can trust, debug, and improve them. The team's work on telemetry for spans, traces, token usage, cost, eval signals, and OpenTelemetry-based normalization across ecosystems like LangChain, OpenLLMetry, and Braintrust is exactly the kind of infrastructure that turns AI from a demo into an operable product.

My recent work is directly relevant. I built Agentic, a non-custodial signing layer for AI agents where the agent can prepare an action, but the user's real wallet remains the approval boundary. It includes MCP tools, Vercel AI tools, a Solana Agent Kit adapter, structured tool schemas, approval polling, transaction and message signing, and a shared WalletBackend abstraction across agent surfaces. That forced me to think carefully about intent, approval, failure states, auditability, and how to make agent behavior inspectable when real user assets are involved.

I also built SolPulse, a live Solana trading platform with 350+ users across web, iOS, Android, and the Solana dApp Store. It uses real-time execution flows, WebSockets/SSE, PostgreSQL/Redis, dashboards, execution logs, wallet-safe signing, and reliability practices around live capital. That experience gave me a practical view of production ownership: when systems are live, observability is not decoration; it is how you protect users, debug quickly, and keep moving without guessing.

I would bring Honeycomb a useful mix of full-stack product engineering, agent infrastructure experience, TypeScript/React, backend/API judgment, and production reliability instincts. I am especially interested in helping define the source of truth for how AI systems behave in production, because I have felt the cost of opaque agent workflows firsthand. Honeycomb's LLM observability work feels like the right layer of the stack to make production AI safer, clearer, and more useful for engineering teams.

Best,
Mathew Stevens