

Mathew Stevens

Pittsburg, CA | mathewstevens7457@gmail.com

Portfolio: <https://mathewstevens.dev/portfolio/> | GitHub: <https://github.com/mstevens843>

SolPulse: <https://solpulse.trade> | Agentic: <https://agentic-signer.com>

June 17, 2026

Labelbox Hiring Team

Re: Applied Research Engineer, Agents

Dear Hiring Team,

I am excited to apply for Labelbox's Applied Research Engineer, Agents role because agent capability increasingly depends on the quality of the data, benchmarks, traces, and feedback loops behind the model. The role's focus on browser interactions, SWE/code traces, GUI sessions, multi-turn workflows, and agent evaluation is especially interesting to me because I have been building practical agent infrastructure where actions need to be reliable, inspectable, and safe enough for real users.

I built Agentic, a non-custodial signing layer for AI agents on Solana. In Agentic, an agent can prepare an action, but the user's real wallet remains the approval boundary. The system includes MCP tools, Vercel AI tools, a Solana Agent Kit adapter, approval polling, transaction/message signing, and a shared WalletBackend abstraction across agent surfaces. That work forced me to think carefully about tool use, structured action generation, user-in-the-loop approval, failure states, and how to make agent behavior observable enough that users can trust what is happening.

I also built SolPulse, a live Solana trading platform with 350+ users across web, iOS, Android, and the Solana dApp Store. It includes real-time execution, dashboards, WebSockets/SSE, PostgreSQL/Redis, wallet-safe signing, and reliability requirements around live capital. Across both projects, the hard part was not just making something work once. It was designing systems that remain useful under real workflows, real latency, and real consequences.

While my background is more applied systems and product engineering than academic research, that is also the value I would bring to Labelbox: the ability to turn agent ideas into working infrastructure, reason about evaluation from the perspective of product correctness, and collaborate with research and customer teams on practical benchmarks and data programs. I am especially interested in building the tooling and datasets that help agents move from promising demos to dependable systems.

Sincerely,

Mathew Stevens